

# **COL 733: CLOUD COMPUTING TECHNOLOGY**

## **FUNDAMENTALS**

### **Assignment 3**

### **Disk Virtualisation**

#### **Group 11**

##### **Members:**

- Ankit Shubham  
2014CS10158
- Ayush Gupta  
2014CS50281
- Deepak Bansal  
2014CS50435
- Kapil Kumar  
2014CS50736

# **DISK VIRTUALISATION (CONSOLIDATION & PARTITIONING)**

## **I. Read/Write of Blocks**

We have a class *BlockData*, the objects of which represent each blocks. The class *BlockMetaData* represents the meta data associated with each block. The *free* variable captures the information whether a block is free or not. The class *FileSystem* represents the main file system that is based on the idea of virtual disks. *diskA* and *diskB* represent the physical disks with capacity of 200 and 300 block respectively. The file system has an array *blocksMetaData* of 500 size which stores the meta data corresponding to each block.

### **Functions**

- **writeBlock (self, blockNum, writeData)**: writes data (*writeData*) to a particular block number (*blockNum*). It initially performs the valid block number check to verify if block number lies between 1 & 500 else returns an Invalid Block Number error. Secondly, it checks whether the size of the *writeData* exceeds the maximum size limit, *blockSize* (100 bytes here). When we write to a block, we update it's free variable to False and store the data either in *diskA* or *diskB* depending on the block number.
- **readBlock (self,blockNum,readData)**: it reads data from a particular block number (*blockNum*) after performing the valid block number check and verifying whether block actually has any data to be read or is free.

### **Tests**

The function *runTests()* performs multiple write and read tests using the above two functions. We try to write normally to a valid block, overwrite a block, writing to invalid blocks and writing larger than *blockSize* data. Then, we try to read from blocks with different checks and conditions. As verifiable by executing the code, we obtain the expected and explainable results.

## **II. Disk Creation & Deletion**

We incorporate the functionality to create or delete virtual disks over the physical disks and blocks that we have dealt with yet. Now the meta data of the block is expanded to also contain variable *allotted* that captures whether a block has been allocated to a certain disk or not and as well the disk id (*diskID*) of the disk it has been allocated to.

We define the following functions in the class *BlockMetaData*:

- **freeFromDisk()**: It sets a block free and also unallocated to any disk.
- **allotToDisk(disID)**: it allots a particular block to the disk with id as *disID*.

We add a class *DiskData* that stores meta data of disk including following variables:

- **idDisk:** ID of a disk that uniquely identifies it among all other disks.
- **sizeDisk:** stores the size of the disk i.e. number of blocks the disk has.
- **startIndexZeroIndexed:** It stores the starting block number for the disk.

## **Functions**

- **createDisk (diskID, numBlocks)** : It attempts to create a disk with the id as *diskID* and number of blocks as *numBlocks*. In the current scheme, we check if our physical disk has got as many contiguous blocks else we raise an error. The *diskList* variable in the *FileSystem* class is updated to contain this disk's id as well.
- **deleteDisk (id)** : It attempts to delete the disk with the id as *id*. After initial validity checks, it essentially calls the function, *freeFromDisk()* to free the blocks that are associated with this disk.
- **writeDisk (diskId, blockNum, writeData)** : It attempts to write the data, *writeData* in the virtual block number, *blockNum*. Internally, It calculates the physical block number to be actually written.
- **readDisk (diskId,blockNum,readData)** : It reads contents of the block with virtual block number as *blockNum* into *readData*. To read data from the physical blocks, it calculates the physical block number of the block and therefore, calls the *readBlock* function.

## **Tests**

We perform comprehensive tests to check the functioning of the above file system with added support for creating, deleting, reading and writing to disks. The *runBlockTests()* first tests the basic functions, *readBlock* and *writeBlock* to write and read to blocks. Once, we see them working successfully, we then perform the disks test, using the *runDiskTests()* function. It tries to create a normal disk, disk with an existing ID, delete a normal disk, delete an already deleted disk and then reads and writes from a disk. We obtain quite expected and explainable results from the above tests.

## **III. Fragmentation**

The above structure of disk is quite elementary and suffers from growth and fragmentation issues. In this disk scheme, we relax the constraint to have all the blocks of a disk in contiguous storage. Therefore, we store the blocks that this disks possess in a list called, *blockList*. In the file system, we have a queue, *freeBlockList*, that stores the block ids of all the blocks which are unallotted to any disk. With this scheme, we again successfully perform all the relevant tests using the function, *runDiskTests()*.

## IV. Block Replication

We now endeavour to provide fault tolerance and reliability to our storage as it's a very pertinent and important feature of any cloud storage system. We incorporate fault tolerance by the concept of block replication wherein we try to store one or multiple replica(s) of each block. So, in case a block gets corrupt, we can recover the information from one of it's replicas. The class *BlockMetaData* is extended to have two more following variables:

- **errorFlag**: it's a boolean variable and is true if a block gets corrupted.
- **replicationBlock**: it stores the id of the block where it's replica is stored.

When we call the function *freeFromDisk* we set the *replicationBlock* attribute for the block to be None again.

The class *DiskData* is extended now to contain another list, *replicaList* of replicas.

The *writeBlock* and *readBlock* function are modified to not be able to read/write in case the block is corrupted and therefore has *errorFlag* attribute as true.

While reading a block, we simulate random read error with a probability of 10%. We do this by generating a random number between 0 and 99 and if it falls before 10 then we mark the block as corrupt.

The *createDisk* function is modified to actually create a disk of double the given size as we create replicas for each block.

## V. SnapShoting

Snapshoting allows the user to restore our disk to an earlier version, called checkpoint. This is a pretty useful function as we can create multiple checkpoints of a disk at different intervals and then get back to one of them when the need arises. We create the following two functions:

- **createCheckpoint(diskId,checkPtId)**: it creates a checkpoint of the disk with id as *diskId* and the checkpoint number as *checkPtId*. It first performs the basic validity checks that whether there exists a disk with id as *diskId* and whether there already exists a checkpoint with same number. We create a snapshot which is a list of tuple of block data and block meta data. This snapshot is stored in *snapshotList* which is a hashmap (dictionary) of all snapshots indexed by checkpoint numbers. The *snapshotList* is included as part of meta data of the disk i.e. in the class *DiskData*.
- **rollBack(diskId,checkPtId)**: It restores the current state of the disk to an earlier version that is captured in a checkpoint with the id as *checkPtId*. So, we replace all the blocks of the disk block by block along with their data and meta data to those stored in the snapshot.